

# Project 2

● Graded

## Student

Cooper Kennelly

## Total Points

87 / 94 pts

## Question 1

Q1 Data Preprocessing 10 / 10 pts

1.1 Q1 a.i 4 / 4 pts

✓ + 4 pts Correct

+ 2 pts Correct mean

+ 2 pts Correct standard deviation

+ 0 pts Incorrect / Incomplete Response

1.2 Q1 a.ii 4 / 4 pts

✓ + 4 pts Correct explanation

+ 2 pts Attempts explanation but doesn't identify maintaining validity of heldout data.

+ 0 pts Incorrect / Incomplete Response

1.3 Q1 b 2 / 2 pts

✓ + 2 pts Fully Correct

+ 0 pts Incorrect / Incomplete Response

## Question 2

### Q2 Convolutional Neural Networks

27 / 30 pts

2.1	<b>Q2 a</b>	3 / 3 pts
	<ul style="list-style-type: none"><li>✓ + 3 pts Fully Correct</li></ul>	
	<ul style="list-style-type: none"><li>+ 1 pt All convolutional layers correct</li></ul>	
	<ul style="list-style-type: none"><li>+ 0.5 pts At least one convolutional layers correct</li></ul>	
	<ul style="list-style-type: none"><li>+ 1 pt Fully-connected layer correct</li></ul>	
	<ul style="list-style-type: none"><li>+ 1 pt Max Pool layers have no learnable float-valued parameters</li></ul>	
	<ul style="list-style-type: none"><li>+ 0 pts Incorrect / Incomplete Response</li></ul>	
2.2	<b>Q2 f.i</b>	4 / 4 pts
	<ul style="list-style-type: none"><li>✓ + 4 pts Describes 2 or more valid sources</li></ul>	
	<ul style="list-style-type: none"><li>+ 2 pts Describes one valid source</li></ul>	
	<ul style="list-style-type: none"><li>+ 0 pts Incorrect / Incomplete Response</li></ul>	
2.3	<b>Q2 f.ii</b>	4 / 6 pts
	<ul style="list-style-type: none"><li>+ 6 pts Fully Correct</li></ul>	
	<ul style="list-style-type: none"><li>✓ + 1 pt Stopping epoch for patience of 5 plot has lowest validation loss 5 epochs before end</li></ul>	
	<ul style="list-style-type: none"><li>✓ + 1 pt Stopping epoch for patience of 10 plot has lowest validation loss 10 epochs before end</li></ul>	
	<ul style="list-style-type: none"><li>+ 2 pts Correctly discusses how patience 5 or 10 is better</li></ul>	
	<ul style="list-style-type: none"><li>✓ + 2 pts Correctly discusses scenario</li></ul>	
	<ul style="list-style-type: none"><li>+ 0 pts Incorrect / Incomplete Response</li></ul>	
2.4	<b>Q2 f.iii</b>	9 / 9 pts
	<ul style="list-style-type: none"><li>✓ + 9 pts Fully Correct</li></ul>	
	<ul style="list-style-type: none"><li>+ 1 pt Correct new FC layer size</li></ul>	
	<ul style="list-style-type: none"><li>+ 2 pts Correct observation</li></ul>	
	<ul style="list-style-type: none"><li>+ 2 pts Correct observation explanation</li></ul>	
	<ul style="list-style-type: none"><li>+ 2 pts Correct AUROC</li></ul>	
	<ul style="list-style-type: none"><li>+ 2 pts Training AUROC is greater than corresponding validation AUROC</li></ul>	
	<ul style="list-style-type: none"><li>+ 0 pts Incorrect / Incomplete Response</li></ul>	

2.5 Q2 g.i

3 / 3 pts

✓ + 3 pts Fully Correct

+ 1 pt Correct training and validation AUROC

+ 1 pt Correct testing AUROC

+ 1 pt Correct testing and validation accuracy

+ 0 pts Incorrect / Incomplete Response

2.6 Q2 g.ii

2 / 2 pts

✓ + 2 pts Fully Correct

+ 1 pt There is little/no overfitting

+ 1 pt Explanation using training/validation performance

+ 0 pts Incorrect / Incomplete Response

2.7 Q2 g.iii

2 / 3 pts

+ 3 pts Fully Correct

✓ + 1 pt Correct discussion between validation and testing performance

✓ + 1 pt States that validation and test data differ in some way

+ 1 pt Correct explanation

+ 0 pts Incorrect / Incomplete Response

### Question 3

Q3 Visualizing what the CNN has learned

10 / 10 pts

3.1

**Q3 a**

5 / 5 pts

✓ + 5 pts Fully Correct

+ 2 pts Correctly calculates  $\alpha$

+ 1 pt Correctly calculates  $\alpha * A + \alpha * A$

+ 2 pts Correctly calculates L

+ 0 pts Incorrect / Incomplete Response

3.2

**Q3 b**

3 / 3 pts

✓ + 3 pts Fully Correct explanation

+ 1.5 pts Explains other features

+ 0 pts Incorrect / Incomplete Response

3.3

**Q3 c**

2 / 2 pts

✓ + 2 pts Fully Correct

+ 0 pts Incorrect / Incomplete Response

#### Question 4

##### Q4.1 Transfer Learning

12 / 14 pts

###### 4.1 Q4.1 c

2 / 2 pts

✓ + 2 pts Fully Correct

+ 1 pt Correct plot and trends

+ 1 pt Correct epoch selected

+ 0 pts Incorrect / Incomplete Response

###### 4.2 Q4.1 d

4 / 4 pts

✓ + 4 pts Fully Correct

+ 1 pt Correct confusion matrix

+ 1 pt Correct landmark with highest accuracy

+ 1 pt Correct landmark with least accuracy

+ 1 pt Valid analysis/justification

+ 0 pts Incorrect / Incomplete Response

###### 4.3 Q4.1 f

6 / 8 pts

+ 8 pts Fully Correct

✓ + 1 pt Correct plot for No frozen layers

✓ + 1 pt Correct plot for First layer frozen

✓ + 1 pt Correct plot for First two layers frozen

✓ + 1 pt Correct plot for All three layers frozen

✓ + 2 pts Correct explanation for transfer learning

+ 2 pts Correct trend in freezing layers with test AUROC

+ 0 pts Incorrect / Incomplete Response

## Question 5

### Q4.2 Data Augmentation

11 / 11 pts

5.1 **Q4.2 b.iii** 8 / 8 pts

✓ + 8 pts Fully Correct

+ 1 pt Same performance in 2g

+ 1 pt Grayscale (discard original) has smallest difference between validation and test AUROC

+ 2 pts Grayscale (discard original) has best test AUROC

+ 4 pts Correct training plots

+ 0 pts Incorrect / blank

5.2 **Q4.2 c** 3 / 3 pts

✓ + 3 pts Correct

+ 0 pts Incorrect / page not selected

## Question 6

### Q5 Challenge - Report

7 / 9 pts

+ 9 pts Fully Correct

✓ + 1 pt Discuss regularization techniques and justification

+ 1 pt Discuss preprocessing and justification

✓ + 1 pt Discuss model architecture and changes made

✓ + 1 pt Discuss hyperparameter selection

+ 1 pt Discuss transfer learning source and target task and setup

✓ + 1 pt Discuss data augmentation and justification

✓ + 1 pt Discuss metrics for evaluation

✓ + 2 pts Write-up demonstrates effort into experimenting and modifying problem setup

+ 0 pts Minimal effort

+ 0 pts Incorrect / Incomplete Response

## Question 7

### Q6 Code Appendix

10 / 10 pts

✓ + 10 pts Fully Correct

+ 0 pts Incorrect / Incomplete Response

Questions assigned to the following page: [1.1](#), [1.2](#), and [1.3](#)

Question 1

A)

- i) The mean values for the RGB channels, as learned from the entire training partition, are: R channel has a mean of 118.263, G channel has a mean of 118.004, and the B channel has a mean of 118.594. Correspondingly, the standard deviations for these channels are: R channel is 66.307, G channel is 69.039, and B channel is 75.332.
- ii) Extracting the mean and standard deviation from the training set, rather than from the other data partitions, is crucial for a few reasons. First, it prevents data leakage. If we use statistics from the entire dataset, which includes our validation and test sets, our model could inadvertently gain information about these sets, leading to an unrealistically optimistic assessment of its performance. It's imperative that our model remains uninformed about the validation and test datasets to ensure that our evaluations are unbiased. Secondly, this approach ensures our model generalizes well. In real-world applications, when we deploy models, we'll normalize new, unseen data points using statistics from the training data. Hence, our validation and testing during the development should mirror this scenario. Lastly, using the training set's statistics ensures that our normalization process remains consistent. If we were to normalize using statistics from the entire dataset, these numbers would shift each time we split our dataset differently, which would be an unstable and inconsistent approach.

B)



Question assigned to the following page: [2.1](#)

Question 2

A)

Filter Size \* Number of Filters \* Depth and include the Bias per Filter as well as parameters for the fully connected layer

$$5 \times 5 (16 \times 3 + 16 \times 64 + 8 \times 64) + 16 + 64 + 8 + 2 + 64 = 39754$$

B)

C)

D)

E)

Question assigned to the following page: [2.2](#)

F)

- i) Overfitting ~ as the training progresses, the model might start to overfit and rely on specific, chance attributes of the training data more than is actually correct in a theoretical manner, especially if the model is complex with many parameters. This can lead to a scenario where the training loss continues to decrease because the model is fitting more closely to the training data, but the validation loss starts to increase since the model is not generalizing well to new, unseen data. In the provided plots, it can be observed that the training loss consistently decreases over the epochs, but the validation loss decreases initially and then increases. This is a classic scenario for overfitting.

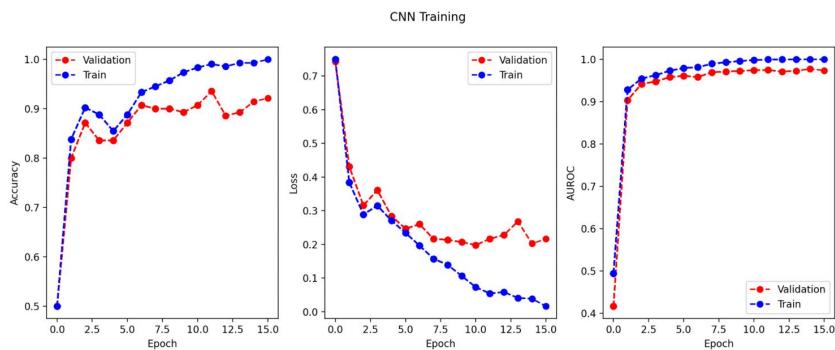
Learning Rate and Optimization Issues ~ the learning rate might be set too high, causing the model to overshoot the optimal values in the loss landscape, leading to fluctuations in the validation loss. Alternatively, if the learning rate is too low, the model might not learn effectively, causing slow or no improvement in the validation loss. It's also possible that the optimizer might get stuck in local minima or saddle points, causing the validation loss not to decrease monotonically.

Questions assigned to the following page: [2.3](#) and [2.4](#)

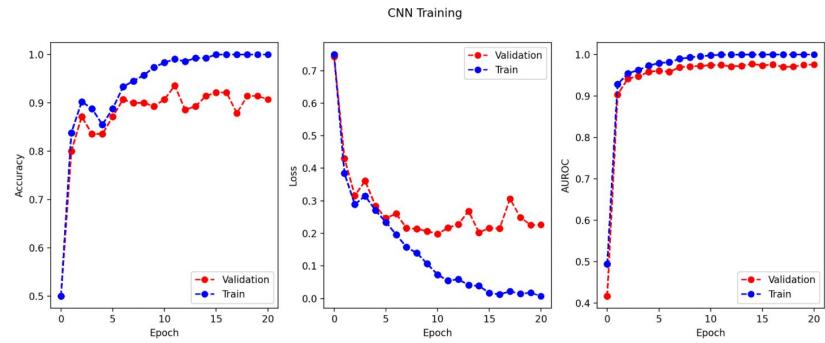
- ii) With a patience of five, it stopped at epoch 15. With a patience of ten, it stops at epoch 20.

Below is the graph of patience = 5

With a patience of ten it stops at Epoch 20 rather than 15 ~ if we evaluate better based on the patience, then both perform equally well since they end based off the same epoch, epoch 10. Increased Patience would work in a data-environment where the training is highly random and where the training hyperplane is highly non-convex, leading to lots of local minima which must be escaped over larger number of epochs to further approximate the global minimum.



Below is the graph of patience = 10



iii)

New Size : 256

Questions assigned to the following page: [2.4](#), [2.5](#), [2.6](#), and [2.7](#)

	EPOCH	Train. AUROC	Val. AUROC
8 filters	10	0.9982	0.9743
64 filters	7	0.9993	0.9694

Differences in plots : The plot for the architecture with 64 filters has training and validation loss diverge much more significantly much sooner ~ it begins to significantly diverge at the third epoch and by the seventh the difference in loss is .2. The plot for the architecture with 8 filters doesn't significantly diverge until the seventh epoch and only gets to a divergence of .2 at the fifteenth training epoch. This suggests that the 64 filter model is too complex and is almost immediately causing overfitting.

G)

i)

	Training	Validation	Testing
Accuracy	0.9833	0.9071	0.625
AUROC	0.9982	0.9743	0.6827

- ii) The difference is not-negligible but also not sufficiently significant to definitively state that there is overtraining. This could be valid performance, we would need to see the trend of validation performance over more training cycles to see if it is increasing. I would say there is not evidence of overtraining present.
- iii) The validation and test performance are significantly different, indicating that there is a large gap between what is represented with our training and validation data, and our test data.

Questions assigned to the following page: [3.1](#), [3.2](#), and [3.3](#)

Question 3

$$A) \alpha_1^1 = \frac{1}{16} \sum \sum \frac{dy^c}{dA_{ij}} = \frac{1}{16} (-1 + 1 - 2 - 1 - 1 + 1 + 1 + 1 + 2 + 2) = \frac{3}{16}$$

$$\alpha_2^1 = \frac{1}{16} \sum \sum \frac{dy^c}{dA_{ij}} = \frac{1}{16} (1 + 2 + 2 + 2 + 2 + 2 + 1 + 1 - 1 - 2 - 1) = \frac{7}{16}$$

$$L^1 = \text{ReLU}(\sum \alpha_i A^1)$$

$$B = [[1, 1, 2, 1], [1, 2, 1, 0], [0, 1, 0, -1], [0, 1, -2, -2]]$$

$$C = [[1, 1, 1, 1], [2, 2, 2, 2], [2, 2, 1, 0], [-1, -1, -1, 0]]$$

$$= \text{ReLU}(\frac{3}{16} * B + \frac{7}{16} * C)$$

$$= [[\frac{5}{8}, \frac{5}{8}, \frac{13}{16}, \frac{5}{8}], [\frac{17}{16}, \frac{5}{4}, \frac{17}{16}, \frac{7}{8}], [\frac{7}{8}, \frac{17}{16}, \frac{7}{16}, \frac{0}{16}], [0, 0, 0, 0]]$$

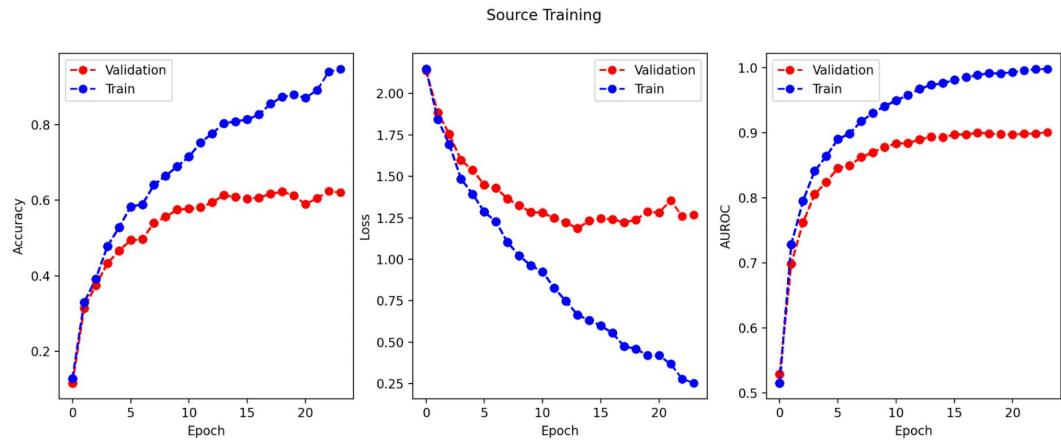
- B) The CNN appears to be using the horse statue as well as the top of the palace to identify the Hofburg Imperial Palace.
- C) Yes, the grad-cam visualizations confirm the performance earlier in the report. The training performance was extremely high, and the differentiation of these two architectures based on the Pantheon's columns and the Hofburg Palace's roofline makes sense as these features are extremely distinguishing between the two of them. However, If the angle of the photos in the test data is such that we do not have a frontal-view of the pantheon or do not have the roof-line of the Hofburg in the photo, then it is likely the model will perform significantly worse.

Question assigned to the following page: [4.1](#)

Question 4

Portion 4.1

- A) (code)
- B) (code)
- C)



Epoch : 13

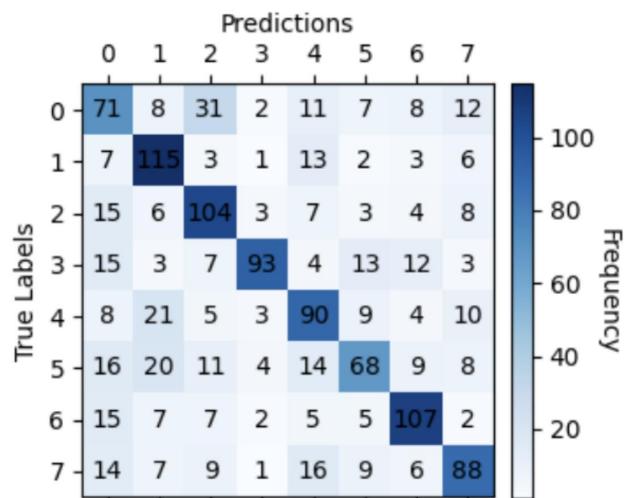
Epoch 13

Validation Accuracy:0.6133  
Validation Loss:1.1857  
Validation AUROC:0.8941  
Train Accuracy:0.8035  
Train Loss:0.6652  
Train AUROC:0.9737

Question assigned to the following page: [4.2](#)

- D) The Petronas Towers, Rialto Bridge, and Hagia Sophia all have extremely distinct geometries and there is only really one perspective on the landmark, so all the images are not only distinct from the other landmarks but also internally very similar leading to a very high true-positive rate.

0 - Colosseum  
 1 - Petronas Towers  
 2 - Rialto Bridge  
 3 - Museu Nacional d'Art de Catalunya  
 4 - St Stephen's Cathedral in Vienna  
 5 - Berlin Cathedral  
 6 - Hagia Sophia  
 7 - Gaudi Casa Batllo in Barcelona



- E) (code)

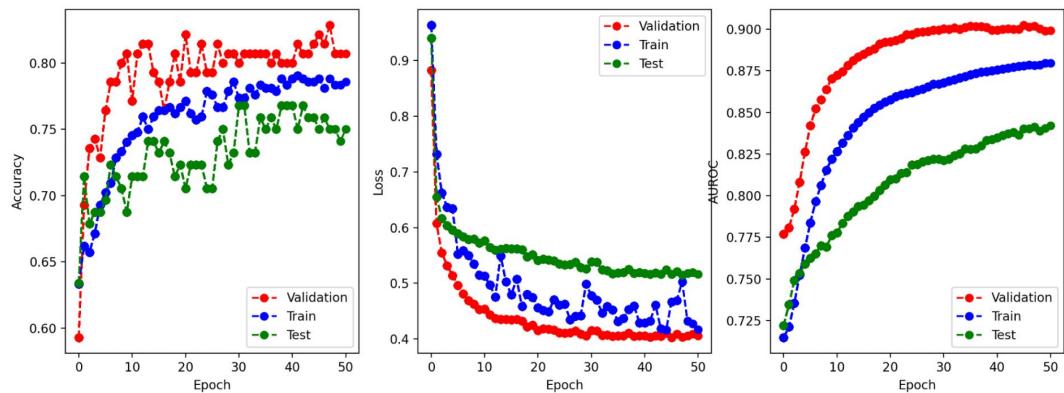
Question assigned to the following page: [4.3](#)

F)

	AUROC TRAIN	AUROC TEST	AUROC VALIDATION
Freeze All Conv	0.878	0.8402	0.9024
Freeze First 2 Conv	0.986	0.8189	0.9624
Freeze First Conv	0.9999	0.809	0.9773
Freeze None	0.983	0.7379	0.9718
No pretraining or Transfer Learning (2(g))	0.9982	0.6827	0.9743

All Convolutional Layers Frozen

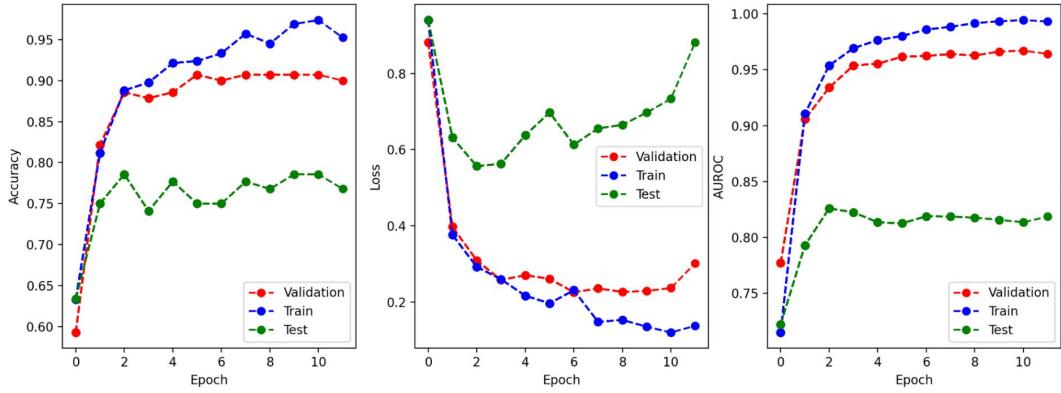
Target Training



Two Frozen Layers

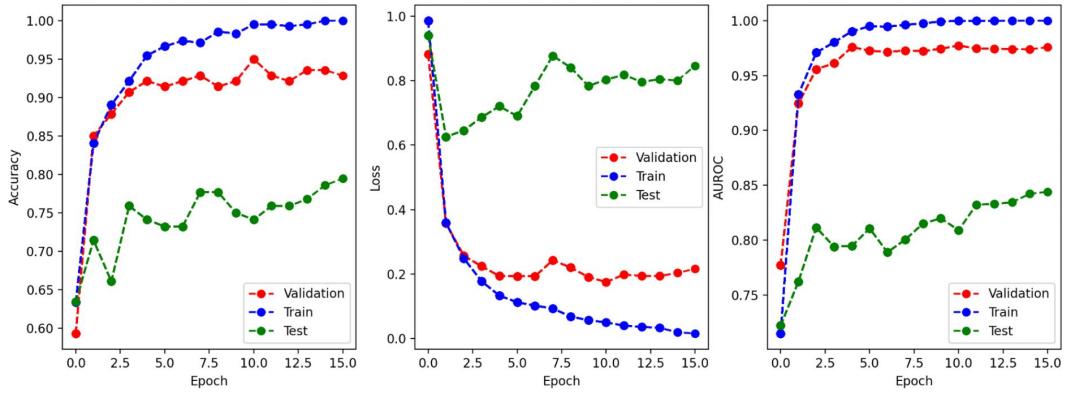
Question assigned to the following page: [4.3](#)

Target Training



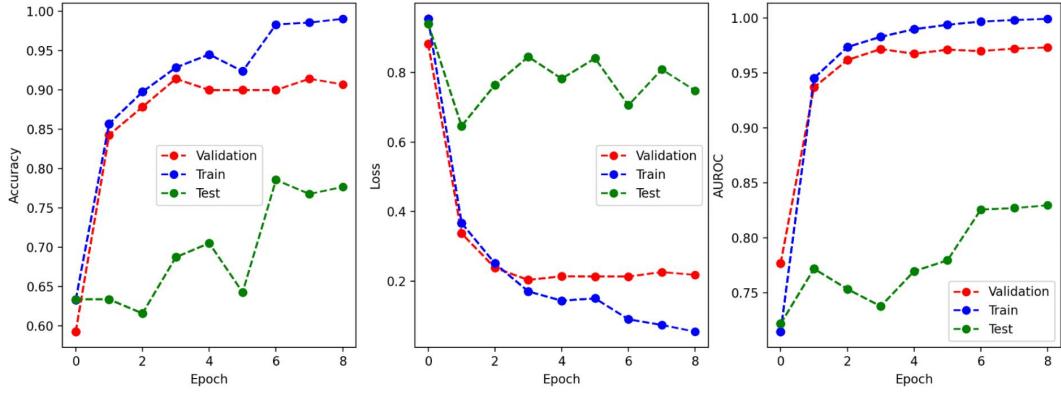
### One Frozen Layer

Target Training



### No Frozen Layers

Target Training



Question assigned to the following page: [4.3](#)

**Effect of Transfer Learning** ~ When we look at the results, using transfer learning seems to have a significant impact on performance. Models that used some form of transfer learning generally outperformed the model that did not leverage source task data, especially in the test set. This suggests that the source task was indeed helpful. Leveraging pre-existing knowledge (weights) from a different but related task can lead to quicker convergence and potentially better generalization to unseen data.

#### Examining Freezing Strategies:

- Freeze All Conv ~ This approach has the lowest train AUROC, indicating that when we freeze all convolutional layers, the model may not be capturing all the intricate patterns from the data. However, its test AUROC is relatively high, which may indicate a good generalization to unseen data.
- Freeze First 2 Conv ~ There's a noticeable drop in test performance compared to freezing all conv layers, suggesting some overfitting might be occurring.
- Freeze First Conv ~ Almost perfect training score suggests potential overfitting, though the validation score remains high. Test performance dropped slightly compared to the previous scenario.
- Freeze None ~ Here, all layers are trainable. This strategy seems to be overfitting the most as seen by a relatively high train score but a much lower test score.

#### Why this might occur

- Freezing layers means the model retains pre-trained features from the source task. When all layers are frozen, the model relies entirely on the source task's features. As layers become trainable, the model adjusts to the target data, which might lead to overfitting if the training data is not representative enough or if the model is too complex.
- Freezing just a subset allows the model to retain some general features (from the frozen layers) while fine-tuning others to the specific task. This can be seen as a balance between leveraging prior knowledge and adapting to new data.
- Not freezing any layer allows complete adaptation to the new task but might also mean losing some useful features from the source task.

**Conclusion** Transfer learning helps in leveraging information from related tasks. However, the extent to which we should adapt these pre-trained models (how many layers we should freeze or make trainable) depends on the similarity between the source and target tasks, as well as the complexity and representativeness of the training data.

Question assigned to the following page: [5.1](#)

Question 4.2

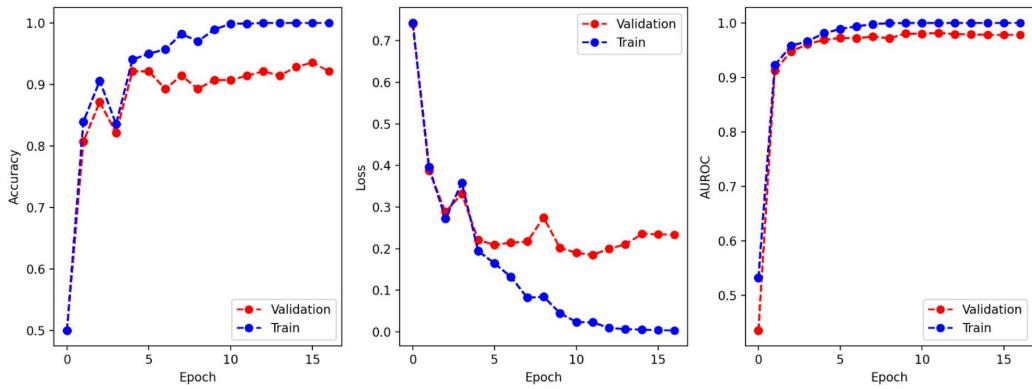
- A) (code)
- B)
- i) (code)
  - ii) (execution of code)
  - iii) The training and validation AUROC are mostly the same across the augmented and regular data except for the Grayscale w/o original Auoc which is significantly lower at .87 rather than .97 for the un-modified data. Furthermore, the test AUROC for all augmentations was higher than for the un-modified data, showing that all the augmentations bolstered model efficacy to some extent. It seems that Grayscale excluding the original photos seemed to have the strongest effect, and I'm not surprised based on the fact that this would likely lend itself best to forcing the model to ascertain geometries and shapes rather than slight / spurious patterns in color and areas.

	AUROC	AUROC	AUROC
	Training	Validation	Test
Rotation (original)(11)	1.0	0.9818	0.751
Grayscale (original)(7)	0.9993	0.9665	0.7554
Grayscale(w/o orig)(7)	0.9889	0.8792	0.8198
No Aug. (2(g))	0.9982	0.9743	0.6827

Rotation With Original Images

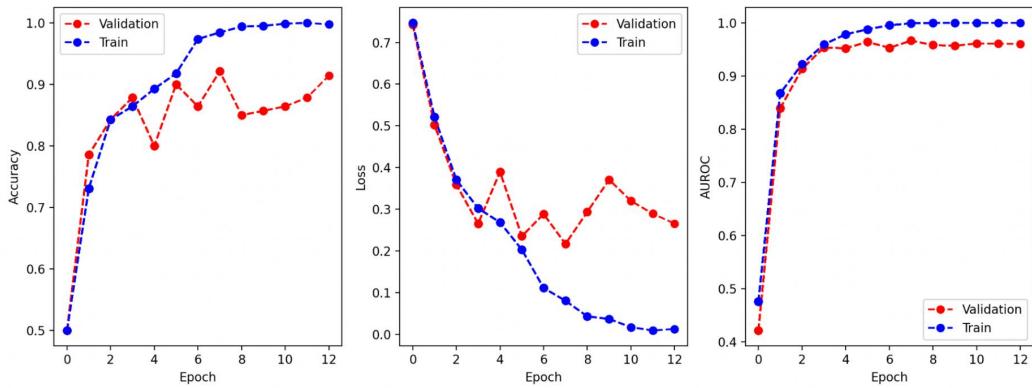
Questions assigned to the following page: [5.1](#) and [5.2](#)

CNN Training



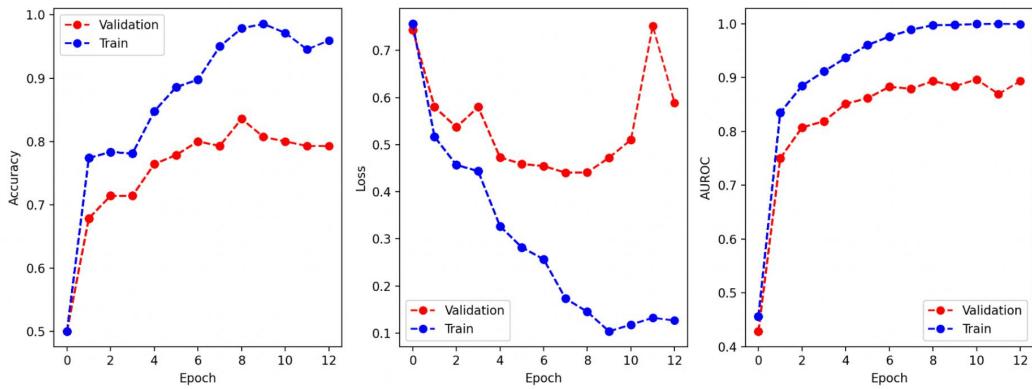
Grayscale With Original Images

CNN Training



Grayscale Without Original Images

CNN Training



C) Upon analyzing the plots, it's clear that the augmentation techniques influence the CNN's training and validation performances differently. When original images are combined with rotation or grayscale augmentations, there's a pronounced

Questions assigned to the following page: [5.2](#) and [6](#)

divergence between training and validation metrics, suggesting overfitting. However, when grayscale augmentation is applied without original images, training and validation metrics are more aligned, indicating reduced overfitting. The change in the plots, as these augmentations are applied, likely stems from how the model interacts with the augmented data. Combining original images with augmentations might be reinforcing certain patterns too strongly, leading the model to over-rely on them, resulting in overfitting. On the other hand, grayscale augmentation without original images may be promoting better generalization. Additionally, the rotation might not provide enough diversity if the rotation range closely resembles the original images, which could explain the overfitting observed in its scenario.

#### Question 5 : Challenge

I will be minimizing for validation loss because I think given that is the grading constraint.

#### Model Architecture :

My first thought for the challenge section was the architecture of the model. People have spent entire PhDs focusing on specific model architectures, so I went online and found a few, and after attempting to implement several of the architectures, AlexNet was the most feasible given our input dimension is non-standard and smaller than all the implemented models. Beyond model architecture, I will also be using augmented data, and assessing performance across different sets of augmented data, specifically rotated with original, grayscale with original, as well as both with original, and both excluding original. After this, I will modify learning rate.

Based on the below tests across learning rate and weight decay for the modified ALexNet architecture, I chose epoch 34, lr=0.00001, weight\_decay=0.005 for the final challenge model. It had the highest validation AUROC across all tested models.

#### AlexNet

For all below : lr=1e-3, weight\_decay=0.001

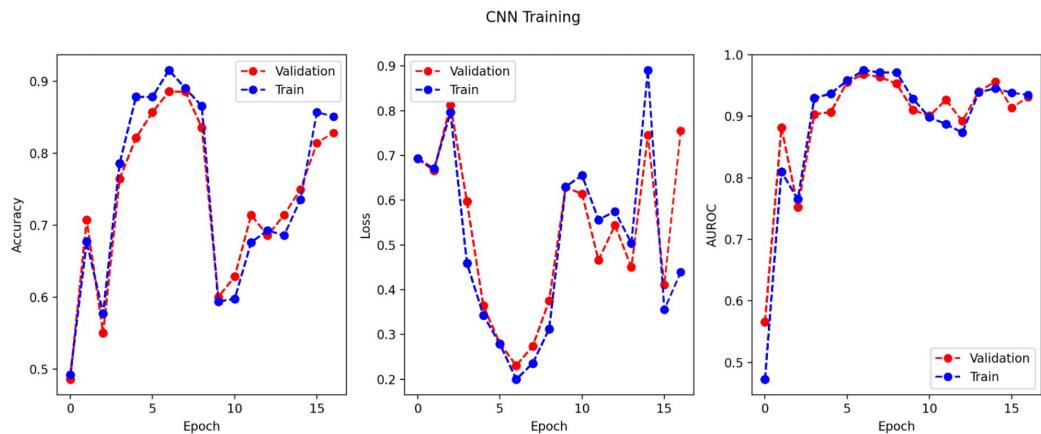
HyperPerams : patience 10, numClassses = 2  
Data : Rotation, Include Original

#### Epoch 6

Validation Accuracy:0.8857  
Validation Loss:0.2306  
Validation AUROC:0.9684  
Train Accuracy:0.9155

Question assigned to the following page: [6](#)

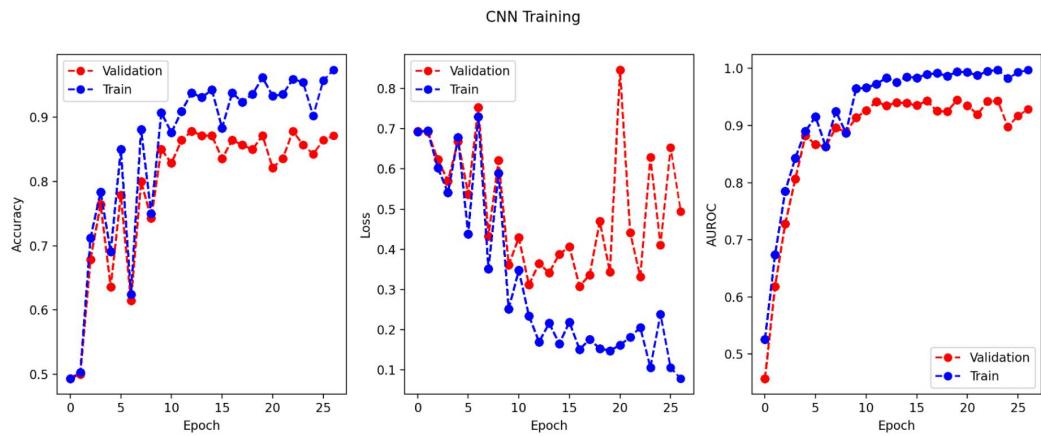
Train Loss:0.1997  
Train AUROC:0.9749



HyperPerams : patience : 10, numClasses = 2  
Data : grayscale, exclude original

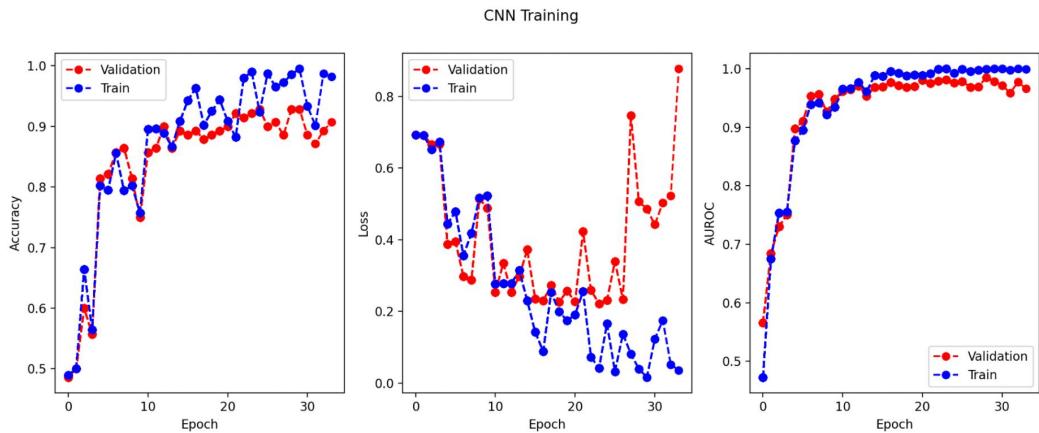
Epoch 16

Validation Accuracy:0.8643  
Validation Loss:0.308  
Validation AUROC:0.9433  
Train Accuracy:0.9381  
Train Loss:0.1508  
Train AUROC:0.9895



HyperPerams : patience 10 , numClasses = 2  
Data, Grayscale, Rotation, Rotation&Grayscale, originals included

Question assigned to the following page: [6](#)



Epoch 23

Validation Accuracy:0.9214  
 Validation Loss:0.2218  
 Validation AUROC:0.981  
 Train Accuracy:0.9905  
 Train Loss:0.0417  
 Train AUROC:0.9997

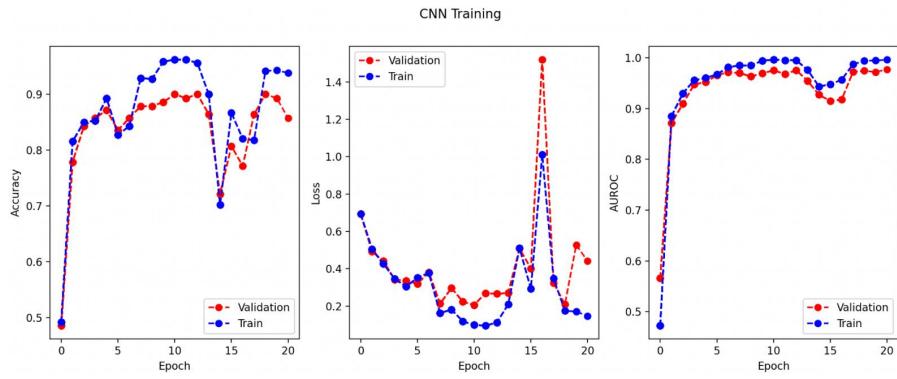
Now that we have determined including the broadest scale of augmented data leads to the highest validation auroc, I will test across a range of learning rates and weight decays

lr=0.001, weight\_decay=0.0005

Epoch 10

Validation Accuracy:0.9  
 Validation Loss:0.2067  
 Validation AUROC:0.9755  
 Train Accuracy:0.9619  
 Train Loss:0.1  
 Train AUROC:0.9965

Question assigned to the following page: [6](#)



$lr=0.01$ ,  $weight\_decay=0.0005$

Epoch 4

Validation Accuracy:0.5214

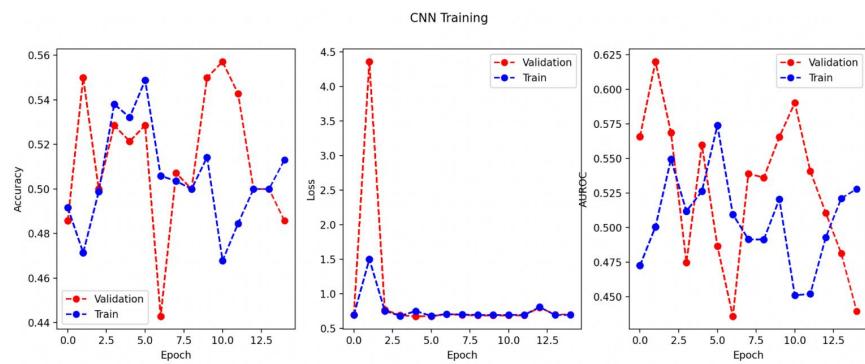
Validation Loss:0.6739

Validation AUROC:0.5598

Train Accuracy:0.5321

Train Loss:0.7501

Train AUROC:0.5263



That was clearly too large based on the initial loss and then massively variant changes in training performance per epoch.

$lr=0.0001$ ,  $weight\_decay=0.0005$

Epoch 9

Validation Accuracy:0.8857

Validation Loss:0.2012

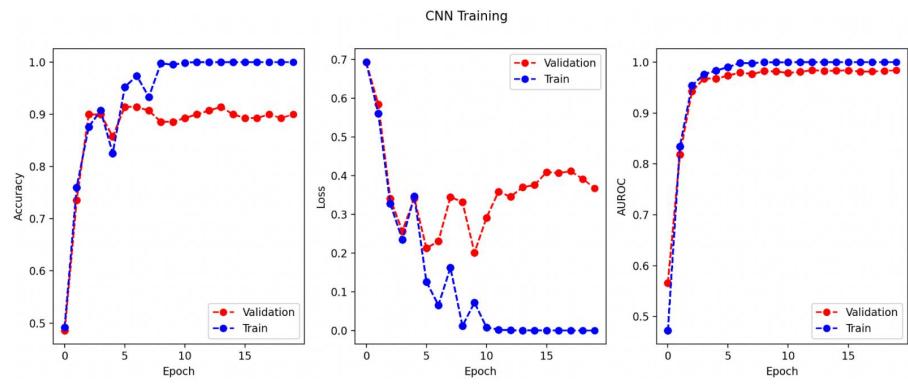
Validation AUROC:0.9816

Train Accuracy:0.9952

Train Loss:0.0727

Train AUROC:0.9998

Question assigned to the following page: [6](#)



Question assigned to the following page: [6](#)

lr=0.00001, weight\_decay=0.0005

Epoch 23

Validation Accuracy:0.9071

Validation Loss:0.1668

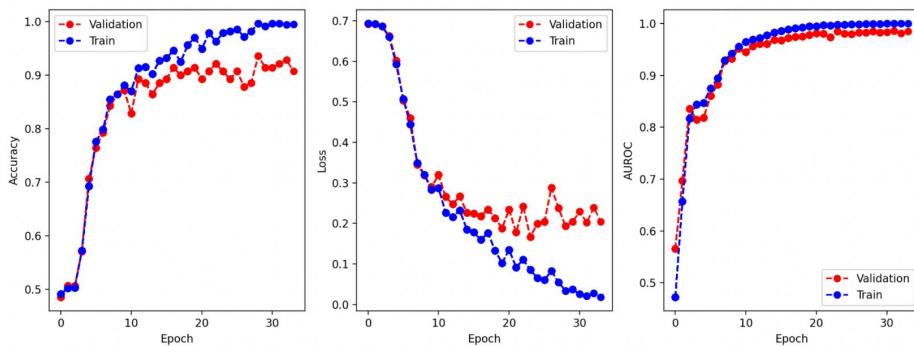
Validation AUROC:0.9845

Train Accuracy:0.9786

Train Loss:0.0853

Train AUROC:0.998

CNN Training



Now lets vary weight decay ...

lr=0.00001, weight\_decay=0.005

Epoch 34

Validation Accuracy:0.9143

Validation Loss:0.1812

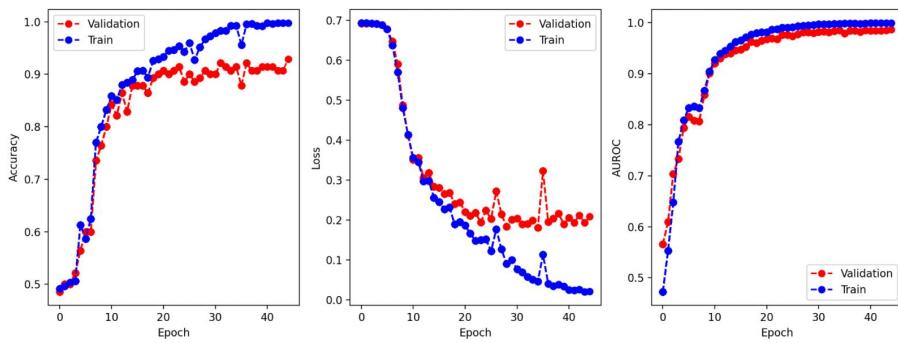
Validation AUROC:0.9853

Train Accuracy:0.9929

Train Loss:0.046

Train AUROC:0.998

CNN Training



lr=0.00001, weight\_decay=0.05

Question assigned to the following page: [6](#)

Epoch 5

Validation Accuracy:0.5286

Validation Loss:0.6927

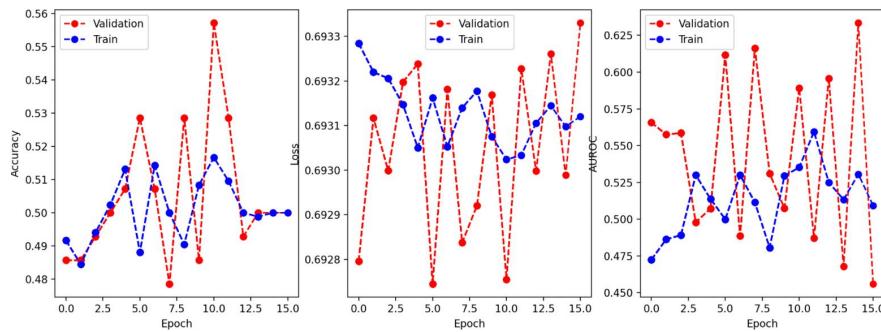
Validation AUROC:0.6116

Train Accuracy:0.4881

Train Loss:0.6932

Train AUROC:0.4999

CNN Training



The decay is definitely too high here based on the per epoch variance in performance ...

lr=0.00001, weight\_decay=0.00005

Epoch 23

Validation Accuracy:0.9143

Validation Loss:0.1744

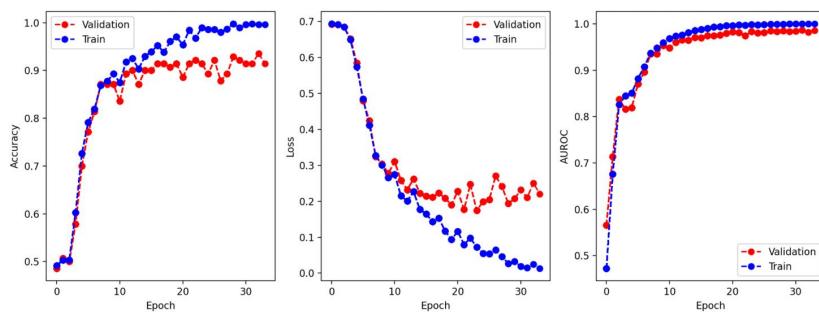
Validation AUROC:0.9835

Train Accuracy:0.9893

Train Loss:0.0722

Train AUROC:0.9986

CNN Training



No questions assigned to the following page.



No questions assigned to the following page.

## CODE APPENDIX

Question assigned to the following page: [7](#)

## Dataset.py

```
1 """
2 EECS 445 - Introduction to Machine Learning
3 Fall 2023 - Project 2
4
5 Landmarks Dataset
6 | Class wrapper for interfacing with the dataset of landmark images
7 | Usage: python dataset.py
8 """
9
10 import os
11 import random
12 import numpy as np
13 import pandas as pd
14 import torch
15 from matplotlib import pyplot as plt
16 from imageio.v3 import imread
17 from PIL import Image
18 from torch.utils.data import Dataset, DataLoader
19 from utils import config
20
21 import img_control
22
23
24 def get_train_val_test_loaders(task, batch_size, **kwargs):
25     """Return DataLoaders for train, val and test splits.
26
27     Any keyword arguments are forwarded to the LandmarksDataset constructor.
28     """
29     tr, va, te, _ = get_train_val_test_datasets(task, **kwargs)
30
31     tr_loader = DataLoader(tr, batch_size=batch_size, shuffle=True)
32     va_loader = DataLoader(va, batch_size=batch_size, shuffle=False)
33     te_loader = DataLoader(te, batch_size=batch_size, shuffle=False)
34
35     return tr_loader, va_loader, te_loader, tr.get_semantic_label
36
37
38 def get_challenge(task, batch_size, **kwargs):
39     """Return DataLoader for challenge dataset.
40
41     Any keyword arguments are forwarded to the LandmarksDataset constructor.
42     """
43     tr = LandmarksDataset("train", task, **kwargs)
44     ch = LandmarksDataset("challenge", task, **kwargs)
45
46     standardizer = ImageStandardizer()
47     standardizer.fit(tr.X)
48     tr.X = standardizer.transform(tr.X)
49     ch.X = standardizer.transform(ch.X)
50
51     tr.X = tr.X.transpose(0, 3, 1, 2)
52     ch.X = ch.X.transpose(0, 3, 1, 2)
53
54     ch_loader = DataLoader(ch, batch_size=batch_size, shuffle=False)
55     return ch_loader, tr.get_semantic_label
56
57
58 def get_train_val_test_datasets(tasks="default", **kwargs):
59     """Return LandmarksDatasets and image standardizer.
60
61     Image standardizer should be fit to train data and applied to all splits.
62     """
63     tr = LandmarksDataset("train", task, **kwargs)
64     va = LandmarksDataset("val", task, **kwargs)
65     te = LandmarksDataset("test", task, **kwargs)
66
67     # Resize
68     # You may want to experiment with resizing images to be smaller
69     # for the challenge portion. How might this affect your training?
70     # tr.X = resize(tr.X)
71     # va.X = resize(va.X)
72     # te.X = resize(te.X)
73
74     # Standardize
75     standardizer = ImageStandardizer()
76     standardizer.fit(tr.X)
77     tr.X = standardizer.transform(tr.X)
78     va.X = standardizer.transform(va.X)
79     te.X = standardizer.transform(te.X)
80
81     # Transpose the dimensions from (N,H,W,C) to (N,C,H,W)
82     tr.X = tr.X.transpose(0, 3, 1, 2)
83     va.X = va.X.transpose(0, 3, 1, 2)
84     te.X = te.X.transpose(0, 3, 1, 2)
85
86     return tr, va, te, standardizer
87
88
89 def resize(X):
90     """Resize the data partition X to the size specified in the config file.
91
92     Use bicubic interpolation for resizing.
93
94     Returns:
95         the resized images as a numpy array.
96     """
97     image_dim = config("image_dim")
98     image_size = (image_dim, image_dim)
99     resized = []
100    for i in range(X.shape[0]):
101        xi = Image.fromarray(X[i]).resize(image_size, resample=2)
102        resized.append(xi)
103    resized = [np.asarray(im) for im in resized]
104    resized = np.array(resized)
105
106    return resized
```

Question assigned to the following page: [7](#)

```

9  class ImageStandardizer(object):
10     """Standardize a batch of images to mean 0 and variance 1.
11
12     The standardization should be applied separately to each channel.
13     The mean and standard deviation parameters are computed in 'fit(X)' and
14     applied using 'transform(X)'.
15
16     X has shape (N, image_height, image_width, color_channel), where N is
17     the number of images in the set.
18
19
20     def __init__(self):
21         """Initialize mean and standard deviations to None."""
22         super().__init__()
23         self.image_mean = None
24         self.image_std = None
25
26     def fit(self, X):
27         """Calculate per-channel mean and standard deviation from dataset X.
28         Hint: you may find the axis parameter helpful"""
29
30         # Compute mean and standard deviation for each channel across all images
31         # X has shape (N, image_height, image_width, color_channel)
32         self.image_mean = np.mean(X, axis=0, 1, 2), dtype=np.float64
33         self.image_std = np.std(X, axis=0, 1, 2), dtype=np.float64
34
35     def transform(self, X):
36         """Return standardized dataset given dataset X."""
37
38         # Check if fit() was called before
39         if self.image_mean is None or self.image_std is None:
40             raise ValueError("Must call fit() before transform()")
41
42         # Subtract the mean and divide by the standard deviation for each channel
43         X_normalized = (X - self.image_mean) / self.image_std
44
45
46 class LandmarksDataset(Dataset):
47     """Dataset class for landmark images."""
48
49     def __init__(self, partition, task="target", augment=False):
50         """Read in the necessary data from disk.
51
52         For parts 2, 3 and data augmentation, 'task' should be "target".
53         For source task of part 4, 'task' should be "source".
54
55         For data augmentation, 'augment' should be True.
56         """
57         super().__init__()
58
59         if partition not in ["train", "val", "test", "challenge"]:
60             raise ValueError("Partition {} does not exist".format(partition))
61
62         np.random.seed(42)
63         torch.manual_seed(42)
64         random.seed(42)
65         self.partition = partition
66         self.task = task
67         self.augment = augment
68
69         # Load in all the data we need from disk
70         if task == "target" or task == "source":
71             self.metadata = pd.read_csv(config["csv_file"])
72
73         if self.augment:
74             print("Augmented")
75             self.metadata = pd.read_csv(config["augmented_csv_file"])
76             self.X, self.y = self._load_data()
77
78         self.semantic_labels = dict(
79             zip(
80                 self.metadata[self.metadata.task == self.task][["numeric_label"]],
81                 self.metadata[self.metadata.task == self.task][["semantic_label"]],
82             )
83         )
84
85     def __len__(self):
86         """Return size of dataset."""
87         return len(self.X)
88
89     def __getitem__(self, idx):
90         """Return (image, label) pair at index 'idx' of dataset."""
91         return torch.from_numpy(self.X[idx]).float(), torch.tensor(self.y[idx]).long()
92
93     def _load_data(self):
94         """Load a single data partition from file."""
95         print("loading %s... %s" % self.partition)
96         df = self.metadata[
97             (self.metadata.task == self.task)
98             & (self.metadata.partition == self.partition)
99         ]
100
101         if self.augment:
102             path = config["augmented_image_path"]
103         else:
104             path = config["image_path"]
105
106         X, y = [], []
107         for i, row in df.iterrows():
108             label = row["numeric_label"]
109             image = imread(os.path.join(path, row["filename"]))
110             X.append(image)
111             y.append(row["numeric_label"])
112
113         return np.array(X), np.array(y)
114
115     def get_semantic_label(self, numeric_label):
116         """Return the string representation of the numeric class label.
117
118         (e.g., the numeric label 1 maps to the semantic label 'hofburg_imperial_palm'
119         """
120         return self.semantic_labels[numeric_label]
121
122
123     if __name__ == "__main__":
124         np.set_printoptions(precision=3)
125         tr, va, te = standardizer.fit_train_val_test_datasets(tasks="farnet", augment=True)

```

Question assigned to the following page: [7](#)

## Target.py

```
1 """
2 EEECS 445 - Introduction to Machine Learning
3 Fall 2023 - Project 2
4 Target CNN
5     Constructs a pytorch model for a convolutional neural network
6     Usage: from model.target import target
7 """
8 import torch
9 import torch.nn as nn
10 import torch.nn.functional as F
11 from math import sqrt
12 from utils import config
13
14 class Target(nn.Module):
15     def __init__(self):
16         """
17             Define the architecture, i.e. what layers our network contains.
18             At the end of __init__() we call init_weights() to initialize all model parameters (weights and biases)
19             in all layers to desired distributions.
20         """
21         super().__init__()
22
23         # Layer 1: Convolutional Layer 1
24         self.conv1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=5, stride=2, padding=2)
25         self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
26
27         # Layer 3: Convolutional Layer 2
28         self.conv2 = nn.Conv2d(in_channels=16, out_channels=64, kernel_size=5, stride=2, padding=2)
29
30         # Layer 5: Convolutional Layer 3
31         self.conv3 = nn.Conv2d(in_channels=64, out_channels=8, kernel_size=5, stride=2, padding=2)
32
33         # Layer 6: Fully connected layer 1 (Output layer)
34         self.fc_1 = nn.Linear(8 * 2 * 2, 2)
35
36         self.init_weights()
37
38     def init_weights(self):
39         """
40             Initialize all model parameters (weights and biases) in all layers to desired distributions"""
41         torch.manual_seed(42)
42
43         for conv in [self.conv1, self.conv2, self.conv3]:
44             C_in = conv.weight.size(1)
45             nn.init.normal_(conv.weight, 0.0, 1 / sqrt(5 * 5 * C_in))
46             nn.init.constant_(conv.bias, 0.0)
47
48         # initialize the parameters for [self.fc_1]
49         nn.init.normal_(self.fc_1.weight, 0.0, 1 / sqrt(self.fc_1.weight.size(1)))
50         nn.init.constant_(self.fc_1.bias, 0.0)
51
52     def forward(self, x):
53         """
54             This function defines the forward propagation for a batch of input examples, by
55             successively passing output of the previous layer as the input into the next layer (after applying
56             activation functions), and returning the final output as a torch.Tensor object.
57         """
58         N, C, H, W = x.shape
59
60         # Forward pass
61         x = F.relu(self.conv1(x)) # Apply conv1 and ReLU
62         x = self.pool(x) # Apply max pooling
63
64         x = F.relu(self.conv2(x)) # Apply conv2 and ReLU
65         x = self.pool(x) # Apply max pooling
66
67         x = F.relu(self.conv3(x)) # Apply conv3 and ReLU
68
69         x = x.view(N, -1) # Flatten the output tensor
70         x = self.fc_1(x) # Apply fully connected layer
71
72         return x
```

Question assigned to the following page: [7](#)

## Train\_common.py 1/2

```
0  from util import config
1  import numpy as np
2  import itertools
3  import os
4  import torch
5  from torch.nn.functional import softmax
6  from sklearn import metrics
7  import utils
8
9
10 def count_parameters(model):
11     """Count number of learnable parameters."""
12     return sum(p.numel() for p in model.parameters() if p.requires_grad)
13
14
15 def save_checkpoint(model, epoch, checkpoint_dir, stats):
16     """Save a checkpoint file to 'checkpoint_dir'."""
17     state = {
18         "epoch": epoch,
19         "state_dict": model.state_dict(),
20         "stats": stats,
21     }
22
23     filename = os.path.join(checkpoint_dir, "epoch{}.checkpoint.pth.tar".format(epoch))
24     torch.save(state, filename)
25
26
27 def check_for_augmented_data(data_dir):
28     """Ask to download augmented data if 'augmented_landmarks.csv' exists in the data directory."""
29     if "augmented_landmarks.csv" in os.listdir(data_dir):
30         print("Augmented data found, would you like to use it? y/n")
31         print("=> ", end="")
32         rep = str(input())
33         return rep == "y"
34     return False
35
36
37 def restore_checkpoint(model, checkpoint_dir, cuda=False, force=False, pretrain=False):
38     """Restore model from checkpoint if it exists.
39
40     Returns the model and the current epoch.
41
42     try:
43         cp_files = [
44             file_
45             for file_ in os.listdir(checkpoint_dir)
46             if file_.startswith("epoch") and file_.endswith(".checkpoint.pth.tar")]
47     except FileNotFoundError:
48         cp_files = None
49     os.makedirs(checkpoint_dir)
50     if not cp_files:
51         print("No saved model parameters found")
52         if force:
53             raise Exception("Checkpoint not found")
54         else:
55             return model, []
56
57     # Find latest epoch
58     for i in itertools.count(1):
59         if "epoch{}.checkpoint.pth.tar".format(i) in cp_files:
60             epoch = i
61             break
62
63     if not force:
64         print(
65             "Which epoch to load from? Choose in range [0, {}].format(epoch),\n"
66             "Enter 0 to train from scratch.",
67         )
68         print("=> ", end="")
69         inp_epoch = int(input())
70         if inp_epoch not in range(epoch + 1):
71             raise Exception("Invalid epoch number")
72         if inp_epoch == 0:
73             print("Checkpoint not loaded")
74             clear_checkpoint(checkpoint_dir)
75             return model, []
76     else:
77         print("Which epoch to load from? Choose in range [1, {}].format(epoch)")
78         inp_epoch = int(input())
79         if inp_epoch not in range(1, epoch + 1):
80             raise Exception("Invalid epoch number")
81
82     filename = os.path.join(
83         checkpoint_dir, "epoch{}.checkpoint.pth.tar".format(inp_epoch)
84     )
85
86     print("Loading from checkpoint {}.".format(filename))
87
88     if cuda:
89         checkpoint = torch.load(filename)
90     else:
91         # Load GPU model on CPU
92         checkpoint = torch.load(filename, map_location=lambda storage, loc: storage)
93
94     try:
95         start_epoch = checkpoint["epoch"]
96         stats = checkpoint["stats"]
97         if pretrain:
98             model.load_state_dict(checkpoint["state_dict"], strict=False)
99         else:
100             model.load_state_dict(checkpoint["state_dict"])
101         print(
102             "=> Successfully restored checkpoint (trained for {} epochs).format(
103                 checkpoint["epoch"])
104         )
105     except:
106         print("=> Checkpoint not successfully restored")
107         raise
108
109     return model, inp_epoch, stats
```

Question assigned to the following page: [7](#)

## Train\_common.py 2/2

```
13 def clear_checkpoint(checkpoint_dir):
14     """Remove checkpoints in `checkpoint_dir`."""
15     filelist = [f for f in os.listdir(checkpoint_dir) if f.endswith(".pth.tar")]
16     for f in filelist:
17         os.remove(os.path.join(checkpoint_dir, f))
18
19     print("Checkpoint successfully removed")
20
21
22 def early_stopping(stats, curr_count_to_patience, global_min_loss):
23     """
24     Implement early stopping based on the validation loss.
25     """
26
27     # Extract the latest validation loss from stats
28     latest_val_loss = stats[-1][1]
29
30     # If the latest validation loss is greater or equal to the global minimum loss
31     if latest_val_loss >= global_min_loss:
32         curr_count_to_patience += 1
33     else:
34         curr_count_to_patience = 0
35         global_min_loss = latest_val_loss
36
37     return curr_count_to_patience, global_min_loss
38
39
40 def evaluate_epoch(
41     axes,
42     tr_loader,
43     val_loader,
44     te_loader,
45     model,
46     criterion,
47     epoch,
48     stats,
49     include_test=False,
50     update_plot=True,
51     multiclass=False,
52 ):
53     """Evaluate the 'model' on the train and validation set."""
54
55     def _get_metrics(loader):
56         y_true, y_pred, y_score = [], [], []
57         correct, total = 0, 0
58         running_loss = 0
59
60         for X, y in loader:
61             with torch.no_grad():
62                 import ipdb; ipdb.set_trace
63                 output = model(X)
64                 predicted = predictions(output.data)
65                 y_true.append(y)
66                 y_pred.append(predicted)
67                 if not multiclass:
68                     y_score.append(softmax(output.data, dim=1)[:, 1])
69                 else:
70                     y_score.append(softmax(output.data, dim=1))
71                 total += y.size(0)
72                 correct += (predicted == y).sum().item()
73             running_loss.append(criterion(output, y).item())
74
75         y_true = torch.cat(y_true)
76         y_pred = torch.cat(y_pred)
77         y_score = torch.cat(y_score)
78         loss = np.mean(running_loss)
79         acc = correct / total
80         if not multiclass:
81             auroc = metrics.roc_auc_score(y_true, y_score)
82         else:
83             auroc = metrics.roc_auc_score(y_true, y_score, multi_class="ovo")
84
85         return acc, loss, auroc
86
87     train_acc, train_loss, train_auc = _get_metrics(tr_loader)
88     val_acc, val_loss, val_auc = _get_metrics(val_loader)
89
90     stats_at_epoch = [
91         val_acc,
92         val_loss,
93         val_auc,
94         train_acc,
95         train_loss,
96         train_auc,
97     ]
98     if include_test:
99         stats_at_epoch += list(_get_metrics(te_loader))
100
101     stats.append(stats_at_epoch)
102     utils.log_training(epoch, stats)
103     if update_plot:
104         utils.update_training_plot(axes, epoch, stats)
105
106
107 def train_epoch(tr_loader, model, criterion, optimizer):
108     """
109     Train the model for one epoch.
110     """
111     model.train() # Set the model to training mode
112
113     for batch_idx, (data, target) in enumerate(tr_loader):
114         optimizer.zero_grad() # Zero out any gradient accumulation
115         output = model(data) # Forward pass
116         loss = criterion(output, target) # Compute loss
117         loss.backward() # Backward pass
118         optimizer.step() # Update model weights
119
120
121 def predictions(logits):
122     """
123     Determine predicted class index given a tensor of logits.
124     Example: Given tensor([0.2, -0.8], [-0.9, -3.1], [0.5, 2.3]), return tensor([0, 0, 1])
125     """
126     Returns:
127         the predicted class output as a PyTorch Tensor
128         """
129     pred = torch.argmax(logits, dim=1)
130
131     return pred
```

Question assigned to the following page: [7](#)

## Train\_cnn.py

```
10 import torch
11 import numpy as np
12 import random
13 from dataset import get_train_val_test_loaders
14 from model.target import Target
15 from train_common import *
16 from utils import config
17 import utils
18 import rng_control
19
20 def main():
21     """Train CNN and show training plots."""
22     # Data loaders
23     if check_for_augmented_data("./data"):
24         tr_loader, va_loader, te_loader, _ = get_train_val_test_loaders(
25             task="target", batch_size=config("target.batch_size"), augment=True
26         )
27     else:
28         tr_loader, va_loader, te_loader, _ = get_train_val_test_loaders(
29             task="target",
30             batch_size=config("target.batch_size"),
31         )
32     # Model
33     model = Target()
34
35     # Define loss function and optimizer. Replace "None" with the appropriate definitions.
36     criterion = torch.nn.CrossEntropyLoss()
37     optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
38
39     print("Number of float-valued parameters:", count_parameters(model))
40
41     # Attempts to restore the latest checkpoint if exists
42     print("Loading cnn...")
43     model, start_epoch, stats = restore_checkpoint(model, config("target.checkpoint"))
44
45     axes = utils.make_training_plot()
46
47     # Evaluate the randomly initialized model
48     evaluate_epoch(
49         axes, tr_loader, va_loader, te_loader, model, criterion, start_epoch, stats
50     )
51
52     # initial val loss for early stopping
53     global_min_loss = stats[0][1]
54
55     # Define patience for early stopping. Replace "None" with the patience value.
56     patience = 5
57     curr_count_to_patience = 0
58
59     # Loop over the entire dataset multiple times
60     epoch = start_epoch
61     while curr_count_to_patience < patience:
62         # Train model
63         train_epoch(tr_loader, model, criterion, optimizer)
64
65         # Evaluate model
66         evaluate_epoch(
67             axes,
68             tr_loader,
69             va_loader,
70             te_loader,
71             model,
72             criterion,
73             epoch + 1,
74             stats,
75             include_test=False,
76         )
77
78         # Save model parameters
79         save_checkpoint(model, epoch + 1, config("target.checkpoint"), stats)
80
81         # update early stopping parameters
82         curr_count_to_patience, global_min_loss = early_stopping(
83             stats, curr_count_to_patience, global_min_loss
84         )
85
86         epoch += 1
87         print("Finished Training")
88         # Save figure and keep plot open
89         utils.save_cnn_training_plot()
90         utils.hold_training_plot()
91
92     if __name__ == "__main__":
93         main()
```

Question assigned to the following page: [7](#)

## source.py

```
93     import torch
94     import torch.nn as nn
95     import torch.nn.functional as F
96     from math import sqrt
97     from utils import config
98
99
100    class Source(nn.Module):
101        def __init__(self):
102            super().__init__()
103
104            # Define Convolutional layers
105            self.conv1 = nn.Conv2d(in_channels=3, out_channels=16, kernel_size=5, stride=2, padding=2) # padding set to 2 to achieve 'SAME' effect
106            self.conv2 = nn.Conv2d(in_channels=16, out_channels=64, kernel_size=5, stride=2, padding=2) # padding set to 2 to achieve 'SAME' effect
107            self.conv3 = nn.Conv2d(in_channels=64, out_channels=8, kernel_size=5, stride=2, padding=2) # padding set to 2 to achieve 'SAME' effect
108
109            # Define Max Pooling layer
110            self.pool = nn.MaxPool2d(kernel_size=2, stride=2, padding=0)
111
112            # Define Fully Connected layer
113            self.fc1 = nn.Linear(8 * 2 * 2, 8) # 8 channels, each of size 2x2
114
115            self.init_weights()
116
117    def init_weights(self):
118        """Initialize all model parameters (weights and biases) in all layers to desired distributions"""
119
120        torch.manual_seed(42)
121        for conv in [self.conv1, self.conv2, self.conv3]:
122            C_in = conv.weight.size(1)
123            nn.init.normal_(conv.weight, 0.0, 1 / sqrt(5 * 5 * C_in))
124            nn.init.constant_(conv.bias, 0.0)
125
126            # Initialize the parameters for self.fc1
127            nn.init.normal_(self.fc1.weight, 0.0, 1 / sqrt(self.fc1.weight.size(1)))
128            nn.init.constant_(self.fc1.bias, 0.0)
129
130
131    def forward(self, x):
132        # Apply the first convolutional layer
133        bsz, _, _, _ = x.size()
134
135        x = self.conv1(x)
136        x = F.relu(x)
137        x = self.pool(x)
138
139        # Apply the second convolutional layer
140        x = self.conv2(x)
141        x = F.relu(x)
142        x = self.pool(x)
143
144        # Apply the third convolutional layer
145        x = self.conv3(x)
146        x = F.relu(x)
147
148        # Flatten the tensor
149        x = x.view(bsz, -1) # 64 x 32
150
151        # Apply the fully connected layer
152        x = self.fc1(x)
153
154
155    return x
156
157
```

Question assigned to the following page: [7](#)

## Train\_source.py

```
1 train_source.py > ...
2
3     # Model
4     model = Source()
5
6     # Criterion (loss function)
7     criterion = nn.CrossEntropyLoss()
8
9     # Optimizer
10    optimizer = optim.Adam(model.parameters(), lr=1e-3, weight_decay=0.01)
11
12    print("Number of float-valued parameters:", count_parameters(model))
13
14    # Attempts to restore the latest checkpoint if exists
15    print("Loading source...")
16    model, start_epoch, stats = restore_checkpoint(model, config("source.checkpoint"))
17
18    axes = utils.make_training_plot("Source Training")
19    #import ipdb; ipdb.set_trace()
20
21    # Evaluate the randomly initialized model
22    evaluate_epoch(
23        axes,
24        tr_loader,
25        va_loader,
26        te_loader,
27        model,
28        criterion,
29        start_epoch,
30        stats,
31        multiclass=True,
32    )
33
34    # Initial val loss for early stopping
35    global_min_loss = stats[0][1]
36
37    # Patience for early stopping
38    patience = 10
39    curr_count_to_patience = 0
40
41    # Loop over the entire dataset multiple times
42    epoch = start_epoch
43    while curr_count_to_patience < patience:
44        # Train model
45        train_epoch(tr_loader, model, criterion, optimizer)
46
47        # Evaluate model
48        evaluate_epoch(
49            axes,
50            tr_loader,
51            va_loader,
52            te_loader,
53            model,
54            criterion,
55            epoch + 1,
56            stats,
57            multiclass=True,
58        )
59
60        # Save model parameters
61        save_checkpoint(model, epoch + 1, config("source.checkpoint"), stats)
62
63        curr_count_to_patience, global_min_loss = early_stopping(
64            stats, curr_count_to_patience, global_min_loss
65        )
66
67        epoch += 1
68
69        # Save figure and keep plot open
70        print("Finished Training")
71        utils.save_source_training_plot()
72        utils.hold_training_plot()
73
74    if __name__ == "__main__":
75        main()
```

Question assigned to the following page: [7](#)

## Train\_target.py

```
#!/usr/bin/python

# ECCS 440 - Introduction to Machine Learning
# Fall 2023 - Project 2
# Train Target
# Train a convolutional neural network to classify images.
# Periodically output training information, and saves model checkpoints.
# Usage: python train_target.py
# ***

1  num
2  ECCS 440 - Introduction to Machine Learning
3  Fall 2023 - Project 2
4  Train Target
5  # Train a convolutional neural network to classify images.
6  # Periodically output training information, and saves model checkpoints.
7  # Usage: python train_target.py
8  ***
9
10 import torch
11 import numpy as np
12 import os
13 from dataset import get_train_val_test_loaders
14 from model.target import Target
15 from train_common import *
16 from config import config
17 import util
18 import copy
19
20 import ragcontrol
21
22 def freeze_layers(model, num_layers=0):
23     """Stop tracking gradients on selected layers."""
24     layers = model.named_parameters()
25     for i in range(num_layers+1):
26         next(layers).value.requires_grad_(False)
27
28
29
30
31 def train_lr_loader(va_loader, te_loader, model, model_name, num_layers=0):
32     """Train transfer learning model."""
33     # TODO: Define loss function and optimizer. Replace "None" with the appropriate definitions.
34     criterion = torch.nn.CrossEntropyLoss()
35     optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
36
37
38     print("Freezing target model with", num_layers, "layers frozen")
39     model, start_epoch, stats = restore_checkpoint(model, model_name)
40
41     axes = util.make_training_plot("Target Training")
42
43     evaluate_epoch(
44         axes,
45         va_loader,
46         va_loader,
47         te_loader,
48         model,
49         criterion,
50         start_epoch,
51         stats,
52         stats,
53         include_test=True,
54     )
55
56     # Initial val loss for early stopping
57     global_min_loss = stats[0][1]
58
59     # TODO: Define patience for early stopping. Replace "None" with the patience value.
60     patience = 5
61     curr_count_to_patience = 0
62
63     # Loop over the entire dataset multiple times
64     epoch = start_epoch
65     while curr_count_to_patience < patience:
66         # Train model
67         train_lr_loader(model, criterion, optimizer)
68
69         # Evaluate model
70         evaluate_epoch(
71             axes,
72             tr_loader,
73             va_loader,
74             te_loader,
75             model,
76             criterion,
77             epoch + 1,
78             stats,
79             include_test=True,
80         )
81
82         # Save model parameters
83         save_checkpoint(model, epoch + 1, model_name, stats)
84
85         curr_count_to_patience, global_min_loss = early_stopping(
86             stats, curr_count_to_patience, global_min_loss
87         )
88
89     epoch += 1
90
91     print("Finished Training")
92
93     # Keep plot
94     axes.set_xlabel("Training step (num_layers)")
95     util.show_training_plot()
96
97
98 def main():
99     """Train transfer learning model and display training plots.
100
101     Train four different models with (0, 1, 2, 3) layers frozen.
102     """
103
104     # Data loaders
105     tr_loader, va_loader, te_loader, _ = get_train_val_test_loaders()
106     task = "target"
107     batch_size_config["target", "batch_size"] = 1
108
109
110     freeze_none = Target()
111     print("Loading source...")
112     freeze_none = util.restore_checkpoint(freeze_none, "source.checkpoint", force=True, pretrain=True)
113
114     freeze_one = copy.deepcopy(freeze_none)
115     freeze_two = copy.deepcopy(freeze_none)
116     freeze_three = copy.deepcopy(freeze_none)
117
118     freeze_layers(freeze_one, 1)
119     freeze_layers(freeze_two, 2)
120     freeze_layers(freeze_three, 3)
121
122     train_lr_loader(va_loader, te_loader, freeze_none, "./checkpoints/target0/", 0)
123     train_lr_loader(va_loader, te_loader, freeze_one, "./checkpoints/target1/", 1)
124     train_lr_loader(va_loader, te_loader, freeze_two, "./checkpoints/target2/", 2)
125     train_lr_loader(va_loader, te_loader, freeze_three, "./checkpoints/target3/", 3)
126
127
128 if __name__ == "__main__":
129     main()
```

Question assigned to the following page: [7](#)

## Augment\_data.py

```
augment_data.py > ...
8  import argparse
9  import csv
10 import glob
11 import os
12 import sys
13 import tensorflow as tf
14 from skimage import rotate
15 from imgaug import imread, imwrite
16 #
17 import imgaug
18
19 def Rotate(deg=20):
20     """Return function to rotate image."""
21
22     def _rotate(img):
23         """Rotate a random integer amount in the range (-deg, deg) (inclusive).
24
25         Keep the dimensions the same and fill any missing pixels with black.
26
27         img: H x W x C numpy array
28         returns: H x W x C numpy array
29         """
30         # Generate a random integer degree in the range (-deg, deg)
31         rotation_deg = np.random.randint(-deg, deg + 1)
32
33         # Rotate the image and fill the missing pixels with black
34         rotated_img = rotate(img, rotation_deg, reshape=False, mode='constant', cv2=True)
35
36         return rotated_img
37
38     return _rotate
39
40
41 def Grayscale():
42     """Return function to grayscale image."""
43
44     def _grayscale(img):
45         """Return 3-channel grayscale of image.
46
47         Compute grayscale values by taking average across the three channels.
48
49         Round to the nearest integer.
50
51         img: H x W x C numpy array
52         returns: H x W x C numpy array
53         """
54         # Compute the grayscale values
55         grayscale_values = np.mean(img, axis=-1).astype(np.uint8)
56
57         # Convert it back to a 3-channel image
58         grayscale_img = np.stack([grayscale_values] * 3, axis=-1)
59
60         return grayscale_img
61
62     return _grayscale
63
64
65 def augment(filename, transforms, n=1, original=True):
66     """Augment image at filename.
67
68     :param filename: name of image to be augmented
69     :param transforms: list of image transformations
70     :param n: number of augmented images to save
71     :param original: whether to include the original images in the augmented dataset or not
72     :returns: a list of augmented images, where the first image is the original
73
74     """
75     print("Augmenting (%s)" % filename)
76     img = imgaug.imread(filename)
77     res = [img] if original else []
78     for i in range(n):
79         new = img
80         for transform in transforms:
81             new = transform(new)
82         res.append(new)
83     return res
84
85
86 def main(args):
87     """Create augmented dataset."""
88
89     reader = csv.DictReader(open(args.input, "r"), delimiter=",")
90     writer = csv.DictWriter(open(args.output, "w"),
91                           fieldnames=["filename", "semantic_label", "partition", "numeric_label", "task"],
92                           )
93
94     augment_partitions = set(args.partitions)
95
96     # TODO: change 'augmentations' to specify which augmentations to apply
97     augmentations = [Rotate()] # rotation only
98     # augmentations = [Grayscale()] # grayscale only
99     # augmentations = [Rotate(), Grayscale()] # both rotation and grayscale
100
101     writer.writeheader()
102     os.makedirs("%s/augmented" % args.datadir, exist_ok=True)
103     for f in glob.glob("%s/*" % args.datadir):
104         print("Processing (%s)" % f)
105         os.remove(f)
106
107     for row in reader:
108         if row["partition"] not in augment_partitions:
109             continue
110         img = imgaug.imread("%s/images/%s" % (args.datadir, row["filename"]))
111         augmentations, n=1,
112         original=True, # TODO: change to False to exclude original image.
113
114         for i, img in enumerate(imgs):
115             fn = "%s/images/%s-%d.jpg" % (args.datadir, row["filename"], i)
116             writer.writerow({"filename": fn,
117                             "semantic_label": row["semantic_label"],
118                             "partition": row["partition"],
119                             "numeric_label": row["numeric_label"],
120                             "task": row["task"],
121                             })
122
123
124     if __name__ == "__main__":
125         parser = argparse.ArgumentParser()
126         parser.add_argument("input", help="Path to input CSV file")
127         parser.add_argument("output", help="Data directory", default="./data/")
128         parser.add_argument(
129             "-p", "--partitions",
130             nargs="*",
131             help="Partitions [train|val|test|challenge|none] to apply augmentations to. Defaults to train",
132             default="train",
133         )
134         main(parser.parse_args(sys.argv[1:]))
135
136 INSERT--
```

Question assigned to the following page: [7](#)

## Challenge.py

```
model > challenge.py > Challenge > __init__.py > num_classes
1  """
2  EEECS 445 - Introduction to Machine Learning
3  Fall 2023 - Project 2
4  Challenge
5      Constructs a pytorch model for a convolutional neural network
6      Usage: from model.challenge import Challenge
7  """
8  import torch
9  import torch.nn as nn
10 import torch.nn.functional as F
11 from math import sqrt
12 from utils import config
13
14
15 class Challenge(nn.Module):
16     def __init__(self, num_classes=2):
17         super(Challenge, self).__init__()
18         self.features = nn.Sequential(
19             nn.Conv2d(3, 64, kernel_size=11, stride=4, padding=2),
20             nn.ReLU(inplace=True),
21             nn.MaxPool2d(kernel_size=3, stride=2),
22             nn.Conv2d(64, 192, kernel_size=5, padding=2),
23             nn.ReLU(inplace=True),
24             nn.MaxPool2d(kernel_size=3, stride=2),
25             nn.Conv2d(192, 384, kernel_size=3, padding=1),
26             nn.ReLU(inplace=True),
27             nn.Conv2d(384, 256, kernel_size=3, padding=1),
28             nn.ReLU(inplace=True),
29             nn.Conv2d(256, 256, kernel_size=3, padding=1),
30             nn.ReLU(inplace=True),
31             nn.MaxPool2d(kernel_size=3, stride=2),
32         )
33         self.avgpool = nn.AdaptiveAvgPool2d((2, 2))
34         self.classifier = nn.Sequential(
35             nn.Dropout(),
36             nn.Linear(256 * 2 * 2, 4096),
37             nn.ReLU(inplace=True),
38             nn.Dropout(),
39             nn.Linear(4096, 4096),
40             nn.ReLU(inplace=True),
41             nn.Linear(4096, num_classes),
42         )
43
44     def forward(self, x):
45         x = self.features(x)
46         x = self.avgpool(x)
47         x = torch.flatten(x, 1)
48         x = self.classifier(x)
49         return x
50
```